

1. Build the Image classification model by dividing the model into following 4 stages:
 - a. Loading and preprocessing the image data
 - b. Defining the model's architecture
 - c. Training the model
 - d. Estimating the model's performance

```
In [2]: import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt
```

a. Loading and preprocessing the image data

```
In [3]: train_data_dir = 'Datasets/mnist-jpg/train'
test_data_dir = 'Datasets/mnist-jpg/test'

# Image data generator for training data
train_datagen = ImageDataGenerator(
    rescale=1.0/255
)

# Image data generator for testing data
test_datagen = ImageDataGenerator(
    rescale=1.0/255
)

# Create data generators
train_batch_size = 10000
train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(28, 28), # Resize images to 28x28
    batch_size=train_batch_size,
    class_mode='categorical',
    color_mode='grayscale', # Use 'categorical' for one-hot encoded labels
    shuffle=True,
)

# Load test data without labels (class_mode=None)
test_batch_size = 2000
test_generator = test_datagen.flow_from_directory(
    test_data_dir,
    target_size=(28, 28), # Resize images to 28x28
    batch_size=test_batch_size,
    class_mode='categorical', # Use 'categorical' for one-hot encoded labels
    color_mode='grayscale',
    shuffle=True,
)
```

Found 60000 images belonging to 10 classes.
Found 60000 images belonging to 10 classes.

Selecting first batch containing 10000 images

```
In [4]: x_train, y_train = train_generator[0]
        x_test, y_test = test_generator[0]
```

```
In [5]: print(x_train.shape, y_train.shape)

(10000, 28, 28, 1) (10000, 10)
```

b. Defining the model's architecture

```
In [6]: model = models.Sequential()
        model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
        model.add(layers.MaxPooling2D((2, 2)))
        model.add(layers.Flatten())
        model.add(layers.Dense(64, activation='relu'))
        model.add(layers.Dense(10, activation='softmax'))

        model.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

c. Training the model

```
In [8]: model.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))

Epoch 1/5
157/157 [=====] - 2s 11ms/step - loss: 0.5566 - accuracy: 0.8428 - val_loss: 0.2578 - val_accuracy: 0.9315
Epoch 2/5
157/157 [=====] - 1s 10ms/step - loss: 0.2051 - accuracy: 0.9419 - val_loss: 0.1786 - val_accuracy: 0.9500
Epoch 3/5
157/157 [=====] - 2s 10ms/step - loss: 0.1344 - accuracy: 0.9626 - val_loss: 0.1204 - val_accuracy: 0.9730
Epoch 4/5
157/157 [=====] - 2s 10ms/step - loss: 0.0943 - accuracy: 0.9733 - val_loss: 0.1029 - val_accuracy: 0.9725
Epoch 5/5
157/157 [=====] - 2s 10ms/step - loss: 0.0658 - accuracy: 0.9828 - val_loss: 0.0822 - val_accuracy: 0.9795
Out[8]: <keras.src.callbacks.History at 0x7f7b05b74690>
```

d. Estimating the model's performance

```
In [10]: test_loss, test_accuracy = model.evaluate(x_test, y_test)
         print("Loss: ", test_loss)
         print("Accuracy: ", test_accuracy)
```

```
63/63 [=====] - 0s 2ms/step - loss: 0.0822 - accu  
racy: 0.9795  
Loss: 0.08217164129018784  
Accuracy: 0.9794999957084656
```

```
In [12]: n = 30  
plt.imshow(x_test[n])  
predicted_value = model.predict(x_test)  
print("Actual Number: ", np.argmax(y_test[n]))  
print("Predicted Number: ", np.argmax(predicted_value[n]))
```

```
63/63 [=====] - 0s 2ms/step  
Actual Number: 4  
Predicted Number: 4
```

